

Peer Review Form

Reviewer	Rimsha Irfan
Report Reviewed	Jannat Shaheer
Domain / Role / FW	D-04 (AI Credit Scoring) R-B (The Critic) F-I (Utilitarian Ethics)
Course	CS-405 Deep Learning — Spring 2026

1. Framework Application

Jannat actually uses the utilitarian framework rather than just naming it, which is not as common as it should be. The two ethical tensions in Section 3 are the clearest evidence of this. Tension 1 (Section 3.2) connects discriminatory proxies to marginal utility in a specifically utilitarian way: the argument is not just that rural women are excluded, but that they are excluded from the group for whom credit access would produce the highest welfare gain. That is a meaningfully different claim from a generic fairness objection, and it requires the framework to actually do some work. Tension 2 (Section 3.3) extends this by treating unexplained denials as a welfare cost independent of the material credit outcome -- again, that is a consequentialist move, not just a complaint about process.

One place where the analysis stops short: Section 3.4 opens the question of incommensurable harms -- the Sindhi farmer's seasonal loss versus aggregate urban welfare gains -- but then steps back, saying only that "a serious utilitarian analysis arrives at a more sceptical conclusion." That is fair, but it leaves the most interesting question unanswered. If some harms really are incommensurable with aggregate welfare gains, does utilitarianism have anything left to say? Or has the argument at that point drifted into a different framework? This deserved to be pushed further rather than left open.

2. Stakeholder Landscape

Five distinct groups are identified and each is analyzed along the same three axes -- power, risk, information access -- which makes for a genuinely useful map. The treatment of development finance institutions (IFC, ADB) is the strongest part of this section: the observation that they benefit reputationally from financial inclusion narratives while welfare costs fall entirely on borrowers is the kind of uncomfortable point that a lot of reports elide.

One gap: borrowers who receive credit and then face over-indebtedness are not mapped as a distinct stakeholder group. They appear briefly in Section 4.2 (CGAP 2022 data, 18-25% over-indebtedness rates) but not in Section 2. Since the utilitarian critique depends on counting every

welfare effect, this is a real omission -- the over-indebted borrower's situation is quite different from the rejected borrower's, and the report treats them as essentially the same.

3. Pakistan-Specific Context

The legal analysis is more specific than most. PDPA 2023 Section 4(1)(a) and Section 6 are correctly characterized, and the absence of a GDPR Article 22 equivalent is the right legal gap to identify -- it is precisely what makes the critic's argument concrete rather than abstract. SBP CPF 2020 and PTA Act Section 23 are both cited with enough precision to be actually useful.

The cultural observation about shared family phones and wallets registered in husbands' names is the best paragraph in Section 5. That kind of context does not come from academic papers; it requires thinking about how these systems interact with how people in Pakistan actually live.

4. Personal Ethical Positioning

Section 6 does not read like a role description, which is what the rubric is looking for. The distinction Jannat draws between opposing AI credit scoring and opposing this specific deployment model is a real distinction. The note about the word inclusion appearing in every fintech pitch -- with actual outcomes rarely reported at the same prominence -- is a specific and somewhat irritated observation. It sounds like a person who has thought about this, not someone summarizing a position.

5. Constructive Suggestion

The recommendations in Section 7 are well-targeted, but they are not explicitly tied back to the utilitarian tensions developed in Section 3. Recommendation 1 (SBP circular on Urdu explanations) addresses Tension 2 on opacity, and Recommendation 2 (NCAI bias audit) addresses Tension 1 on discriminatory proxies -- but these connections are implicit. Adding one sentence per recommendation that names the specific welfare harm being addressed would make the logical structure much cleaner and would demonstrate that the recommendations emerge from the framework rather than from general intuition.